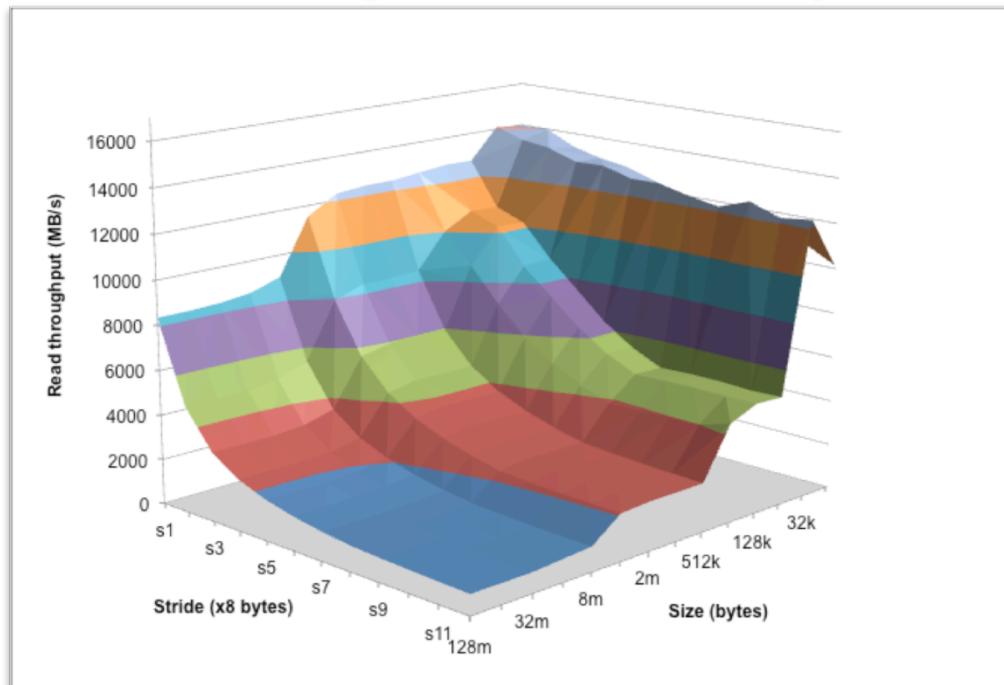


Introducing Computer Systems from a Programmer's Perspective

Randal E. Bryant, David R. O'Hallaron

Computer Science, Electrical & Computer Engineering

Carnegie Mellon University



Outline

Introduction to Computer Systems

- Course taught at CMU since Fall, 1998
- Some ideas on labs, motivations, ...

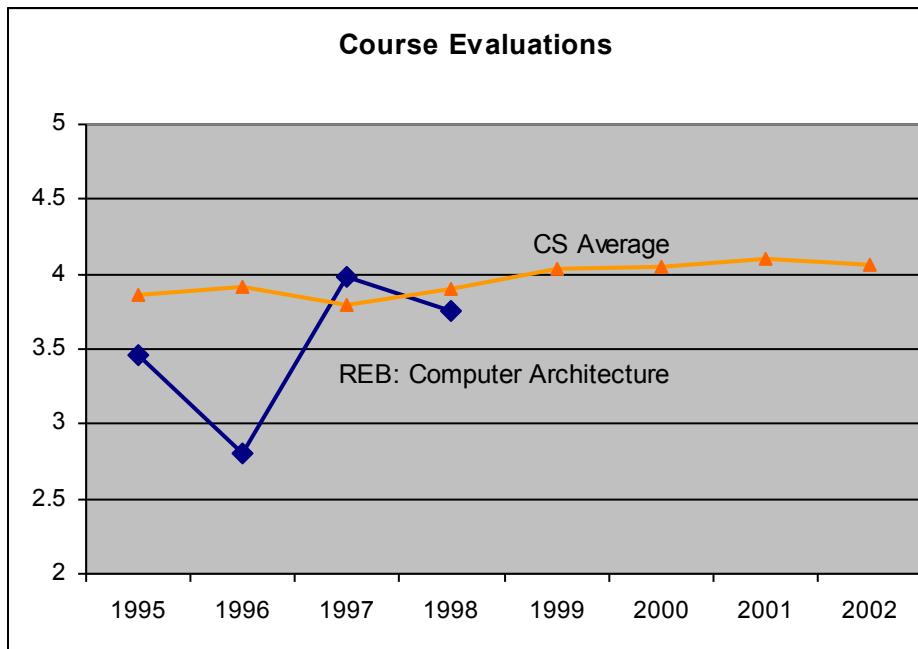
Computer Systems: A Programmer's Perspective

- Our textbook, now in its third edition
- Ways to use the book in different courses

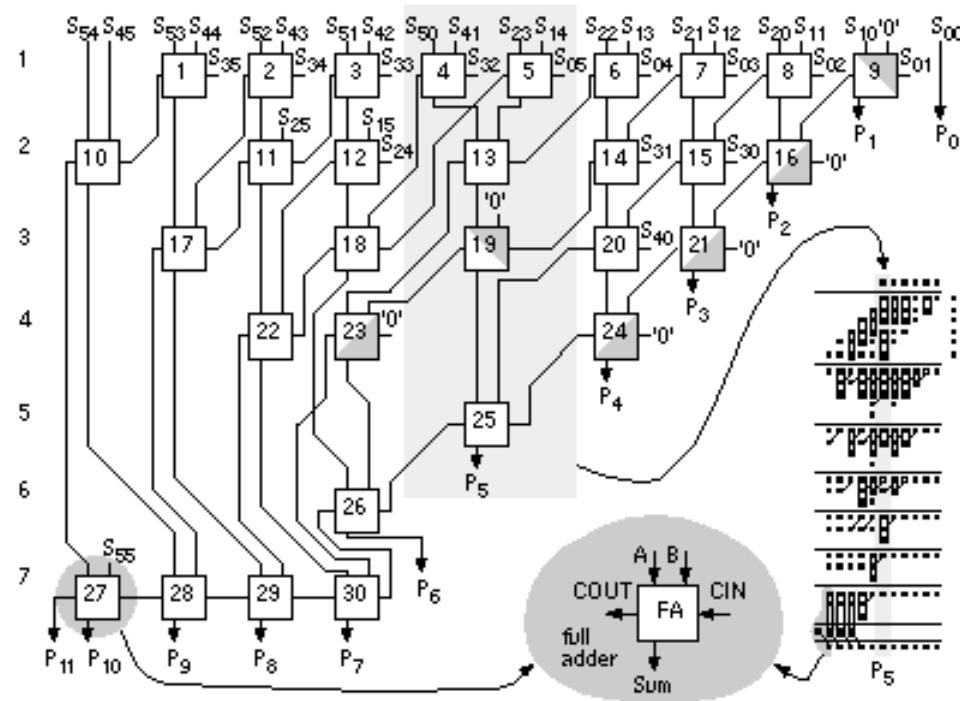
Background

1995-1997: REB/DROH teaching computer architecture course at CMU.

- Good material, dedicated teachers, but students hate it
- Don't see how it will affect their lives as programmers



Computer Arithmetic *Builder's Perspective*



- How to design high performance arithmetic circuits

Computer Arithmetic

Programmer's Perspective

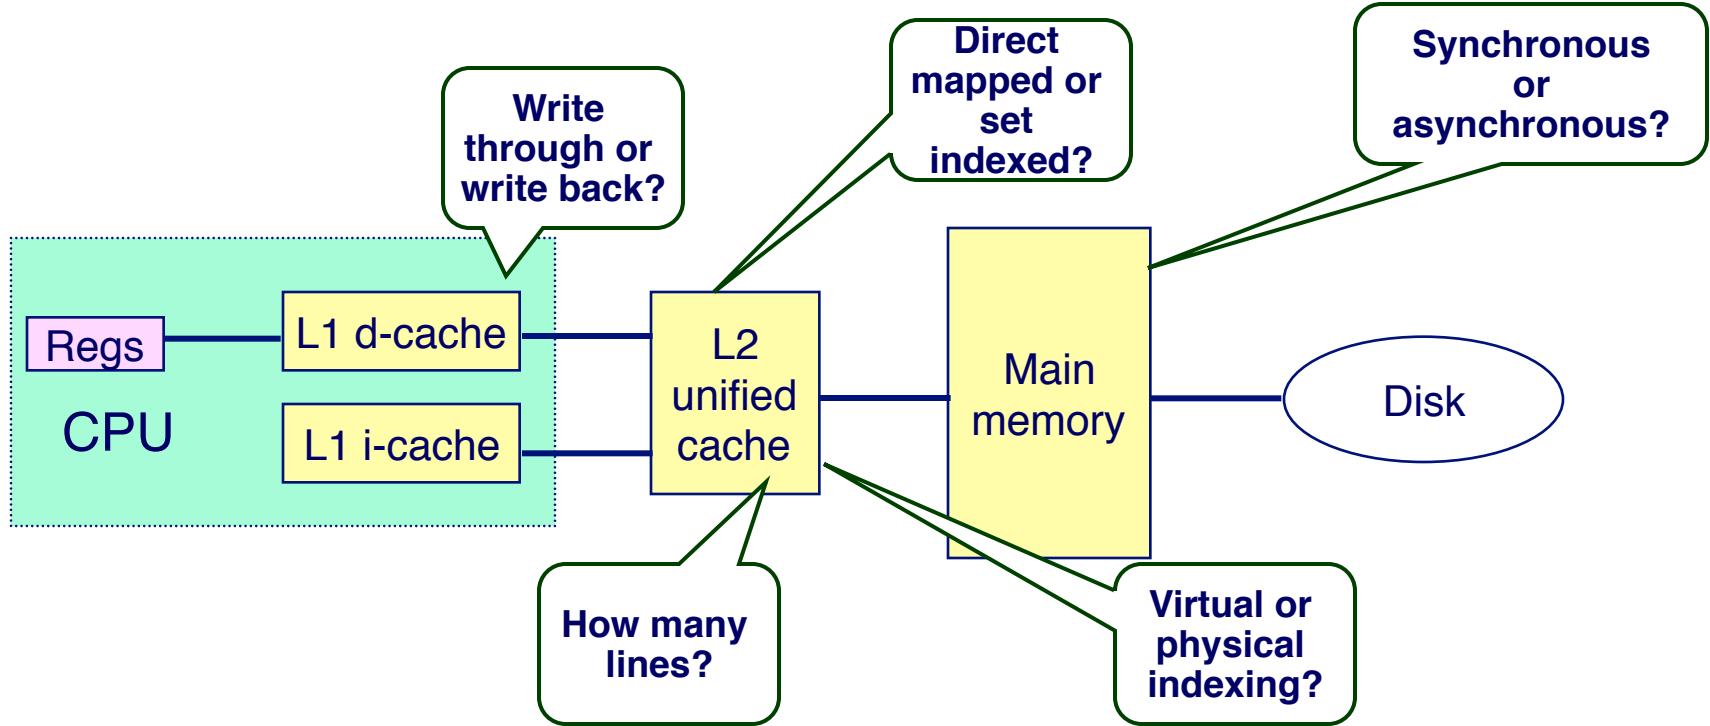
```
void show_squares()
{
    int x;
    for (x = 5; x <= 5000000; x*=10)
        printf("x = %d x^2 = %d\n", x, x*x);
}
```

x =	5	x ² =	25
x =	50	x ² =	2500
x =	500	x ² =	250000
x =	5000	x ² =	25000000
x =	50000	x ² =	-1794967296
x =	500000	x ² =	891896832
x =	5000000	x ² =	-1004630016

- Numbers are represented using a finite word size
- Operations can overflow when values too large
 - But behavior still has clear, mathematical properties

Memory System Builder's Perspective

Builder's Perspective



- Must make many difficult design decisions
- Complex tradeoffs and interactions between components

Memory System Programmer's Perspective

```
void copyij(int src[2048][2048],  
           int dst[2048][2048])  
{  
    int i,j;  
    for (i = 0; i < 2048; i++)  
        for (j = 0; j < 2048; j++)  
            dst[i][j] = src[i][j];  
}
```

4.3 ms

```
void copyji(int src[2048][2048],  
           int dst[2048][2048])  
{  
    int i,j;  
    for (j = 0; j < 2048; j++)  
        for (i = 0; i < 2048; i++)  
            dst[i][j] = src[i][j];  
}
```

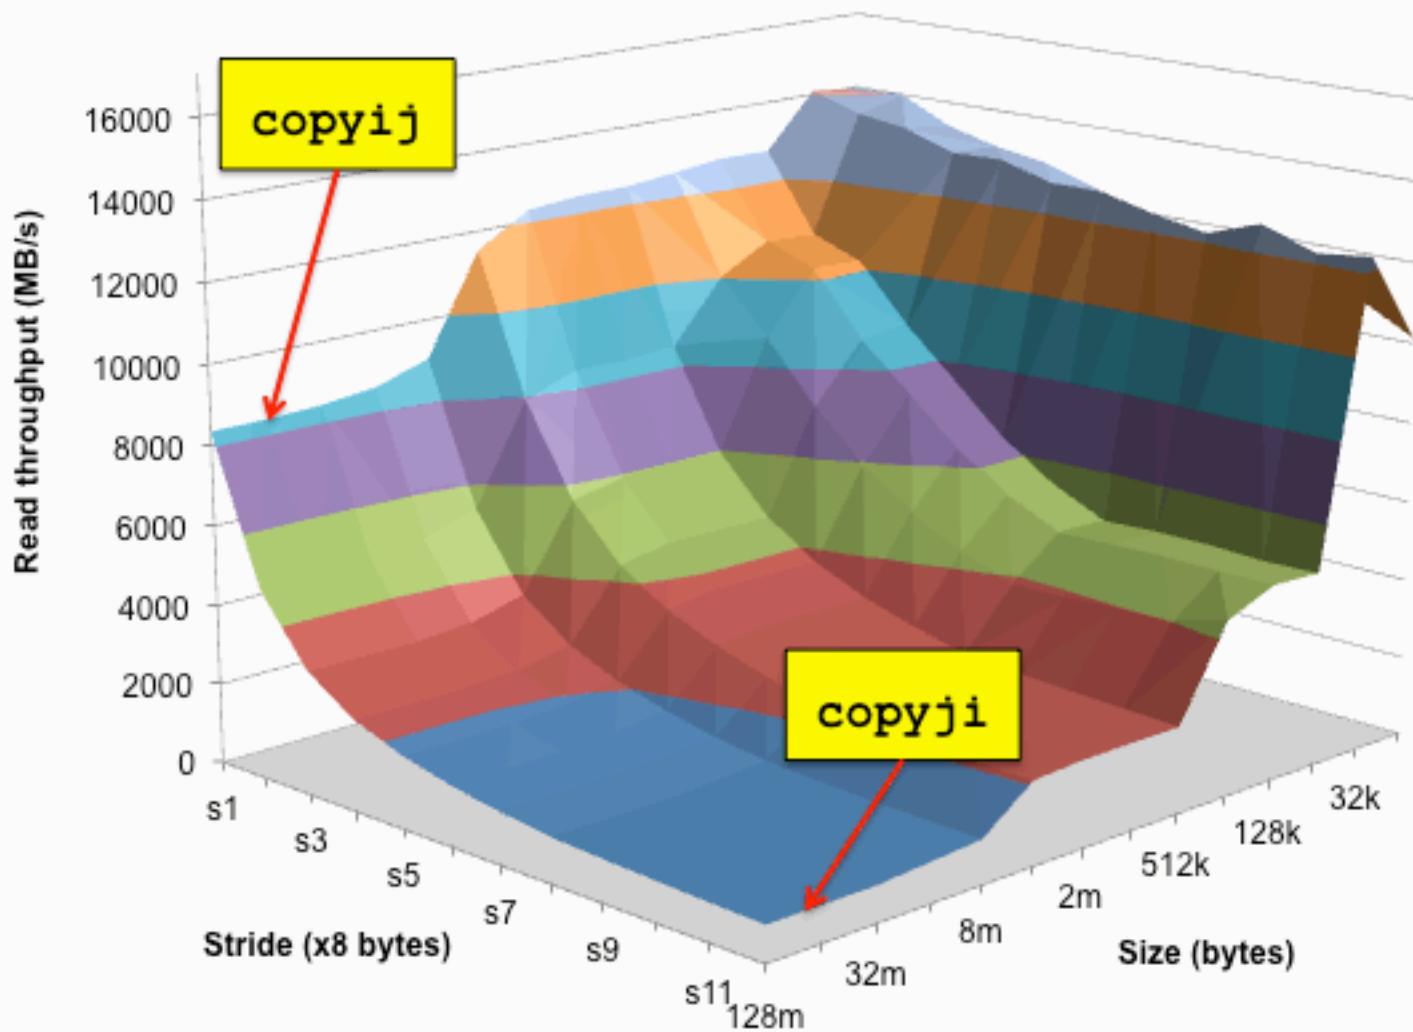
81.8 ms

19 times slower!

(Measured on 2 GHz
Intel Core i7 Haswell)

- Hierarchical memory organization
- Performance depends on access patterns
 - Including how step through multi-dimensional array

The Memory Mountain



Background (Cont.)

1997: OS instructors complain about lack of preparation

- Students don't know machine-level programming well enough
 - What does it mean to store the processor state on the run-time stack?
- Our architecture course was not part of prerequisite stream

Birth of ICS

1997: REB/DROH pursue new idea:

- Introduce them to computer systems from a programmer's perspective rather than a system designer's perspective.
- Topic Filter: What parts of a computer system affect the correctness, performance, and utility of my C programs?

1998: Replace architecture course with new course:

- 15-213: Introduction to Computer Systems

Curriculum Changes

- Sophomore level course
- Eliminated digital design & architecture as required courses for CS majors

15-213: Intro to Computer Systems

Goals

- Teach students to be sophisticated application programmers
 - Immediate value, even if never take another systems course
- Prepare students for upper-level systems courses

Taught every semester to 400+ students

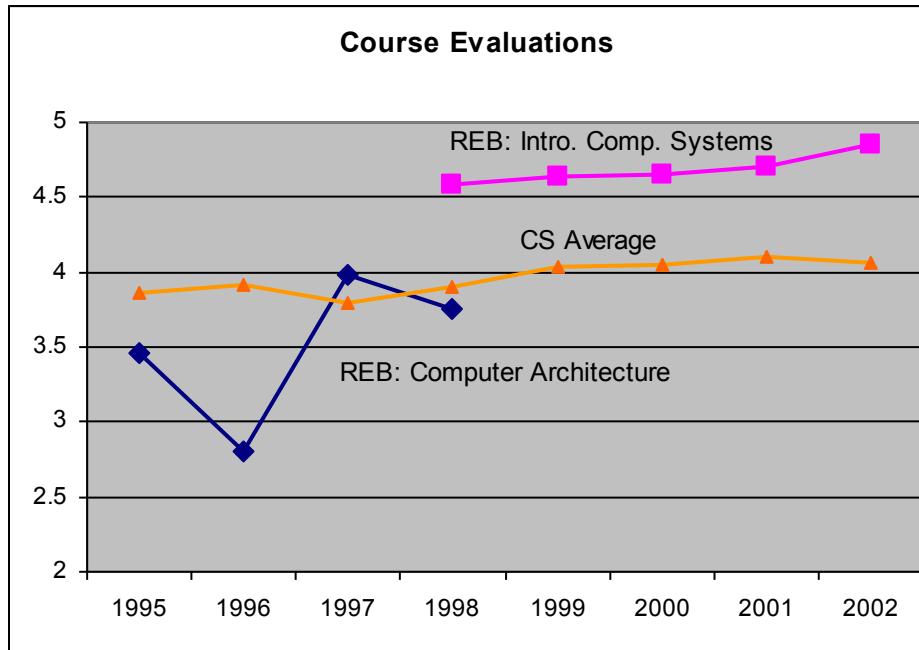
- All CS undergrads (core course)
- All ECE undergrads (core course)
- Many masters students
 - To prepare them for upper-level systems courses
- Variety of others from math, physics, statistics, ...

Preparation

- Optional: Introduction to CS in Python or Ruby
- Imperative programming in C subset

ICS Feedback

Students



Faculty

- Prerequisite for most upper level CS systems courses
- Also required for ECE embedded systems, architecture, and network courses

Lecture Coverage

Data representations [3]

- It's all just bits.
- `int`'s are not integers and `float`'s are not reals.

x86-64 machine language [5]

- Analyzing and understanding compiler-generated machine code.

Program optimization [2]

- Understanding compilers and modern processors.

Memory Hierarchy [3]

- Caches matter!

Linking [1]

- With DLL's, linking is cool again!

Lecture Coverage (cont)

Exceptional Control Flow [2]

- The system includes an operating system that you must interact with.

Virtual memory [4]

- How it works, how to use it, and how to manage it.

Application level concurrency [3]

- Processes and threads
- Races, synchronization

I/O and network programming [4]

- Programs often need to talk to other programs.

Total: 27 lectures, 14 week semester

Labs

Key teaching insight:

- Cool Labs ⇒ Great Course

A set of 1 and 2 week labs define the course.

Guiding principles:

- Be hands on, practical, and fun.
- Be interactive, with continuous feedback from automatic graders
- Find ways to challenge the best while providing worthwhile experience for the rest
- Use healthy competition to maintain high energy.

Lab Exercises

Data Lab (2 weeks)

- Manipulating bits.

Bomb Lab (2 weeks)

- Defusing a binary bomb.

Attack Lab (1 week)

- Buffer overflow and return-oriented programming exploits

Cache Lab (2 weeks)

- Write basic cache simulator and then optimize application

Shell Lab (1 week)

- Writing your own shell with job control.

Malloc Lab (2-3 weeks)

- Writing your own malloc package.

Proxy Lab (2 weeks)

- Writing your own concurrent Web proxy.

Data Lab

Goal: Solve some “bit puzzles” in C using a limited set of logical and arithmetic operators.

- Examples: `absval(x)` , `greaterthan(x,y)` , `log2(x)`

Lessons:

- Information is just bits in context.
- C int's are not the same as integers.
- C float's are not the same as reals.

Infrastructure

- Configurable source-to-source C compiler that checks for compliance.
- Instructor can automatically select from 45 puzzles.
- Automatic testing using formal verification tools

Let's Solve a Bit Puzzle!

```
/*
 * abs - absolute value of x (except returns TMin for TMin)
 * Example: abs(-1) = 1.
 * Legal ops: ! ~ & ^ | + << >>
 * Max ops: 10
 * Rating: 4
 */
int abs(int x) {
    int mask = x>>31;
    return (x^mask) + 1+~mask;
}
```

$$\begin{aligned} 11\dots1_2, &= -1, & x < 0 \\ 00\dots0_2, &= 0, & x \geq 0 \end{aligned}$$

$$\begin{cases} -x - 1, & x < 0 \\ x, & x \geq 0 \end{cases}$$

+

$$\begin{cases} 1, & x < 0 \\ 0, & x \geq 0 \end{cases}$$

=

$$\begin{cases} -x, & x < 0 \\ x, & x \geq 0 \end{cases}$$

Verifying Solutions

```
int abs(int x) {  
    int mask = x>>31;  
    return (x ^ mask) + ~mask + 1;  
}
```

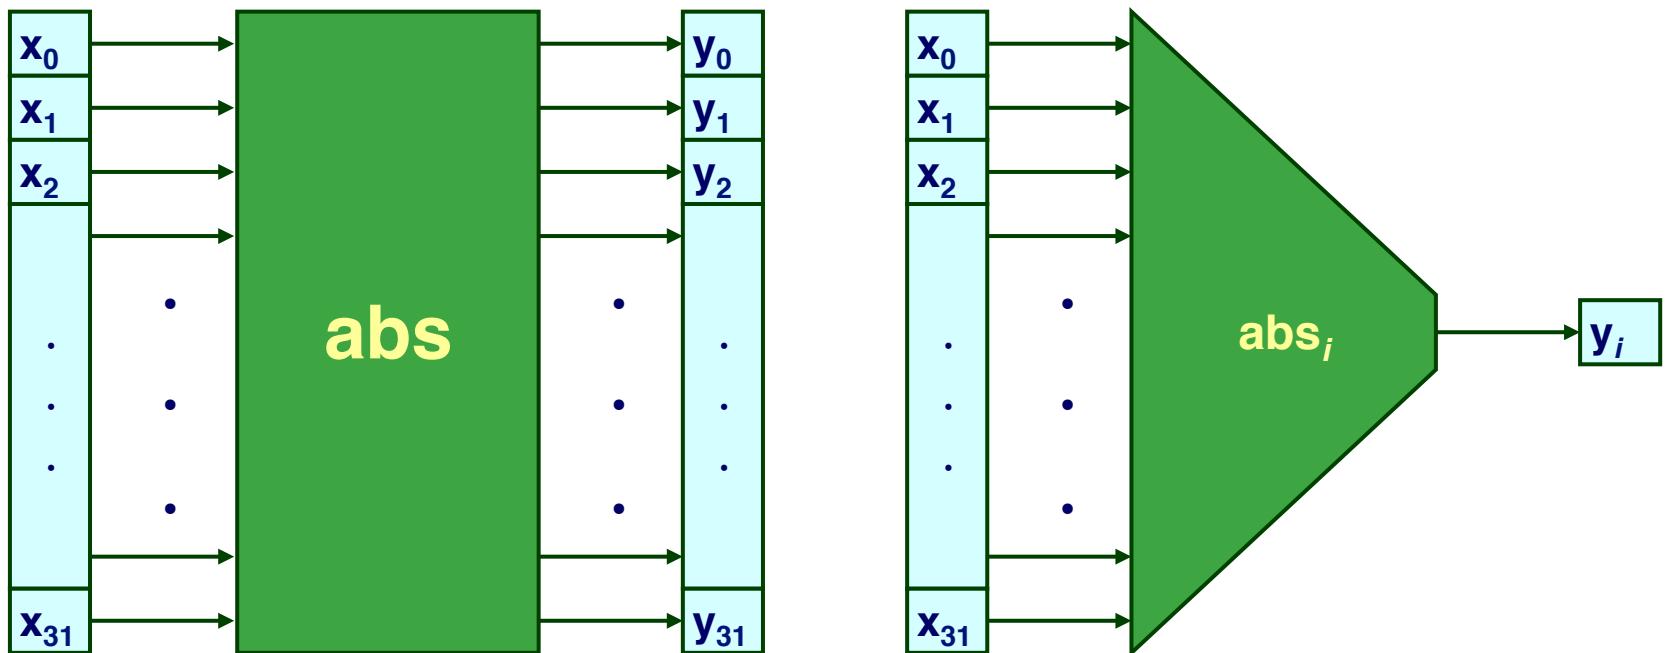
```
int test_abs(int x) {  
    return (x < 0) ? -x : x;  
}
```

Do these functions produce identical results?

How could you find out?

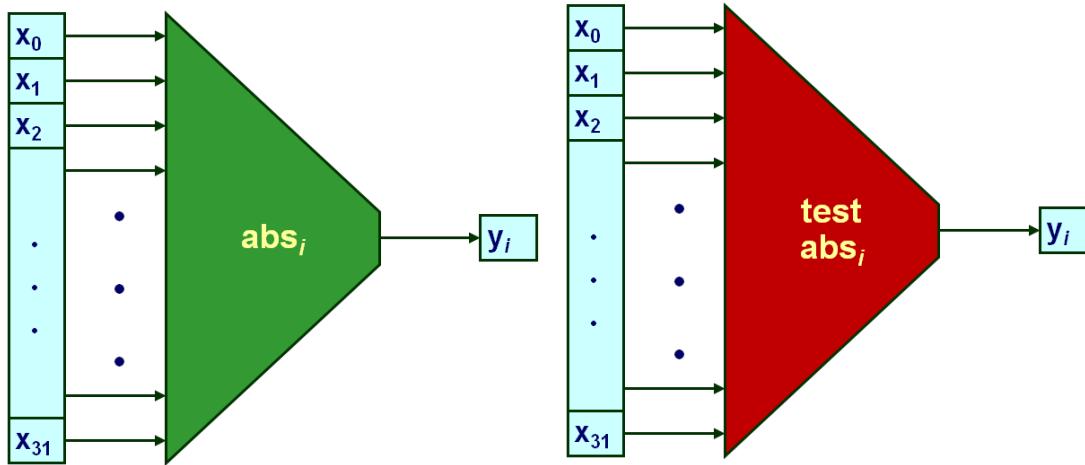
Bit-Level Program Model

```
int abs(int x) {  
    int mask = x>>31;  
    return (x ^ mask) + ~mask + 1;  
}
```



- View computer word as 32 separate bit values
- Each output becomes Boolean function of inputs

Bit-Level Program Verification



- Determine whether functions equivalent for all outputs j
- Exhaustive checking:
 - Single input: $\frac{2^{32} \text{ cases} \times 50 \text{ cycles}}{2 \times 10^9 \text{ cycles / second}} \approx 60 \text{ seconds}$
 - Two input: $2^{64} \text{ cases} \rightarrow 8,800 \text{ years!}$
- Other approaches
 - BDDs, SAT solvers
 - Easily handle these functions (< 1.0 seconds)

Verification Example

```
int iabs(int x) {  
    if (x == 1234567) x++;  
    int mask = x>>31;  
    return (x ^ mask) + ~mask + 1;  
}
```

Almost Correct

- Valid for all but one input value
- Overlooked by our test suite

Counterexample Generation

```
int iabs(int x) {  
    if (x == 1234567) x++;  
    int mask = x>>31;  
    return (x ^ mask) + ~mask + 1;  
}
```

Detected By Checking Code

- Since covers *all* cases
- Generate counterexample to demonstrate problem

```
int main()  
{  
    int val1 = iabs(1234567);  
    int val2 = test_iabs(1234567);  
    printf("iabs(1234567) --> %d [0x%x]\n", val1, val1);  
    printf("test_iabs(1234567) --> %d [0x%x]\n", val2, val2);  
    if (val1 == val2) {  
        printf(.. False negative\n");  
    } else  
        printf(.. A genuine counterexample\n");  
}
```

Bomb Lab

- Idea due to Chris Colohan, TA during inaugural offering

Bomb: C program with six phases.

Each phase expects student to type a specific string.

- Wrong string: bomb *explodes* by printing BOOM! (- ½ pt)
- Correct string: phase *defused* (+10 pts)
- In either case, bomb sends message to grading server
- Server posts current scores anonymously and in real time on Web page

Goal: Defuse the bomb by defusing all six phases.

- For fun, we include an unadvertised seventh *secret phase*

The challenge:

- Each student get only binary executable of a *unique* bomb
- To defuse their bomb, students must disassemble and reverse engineer this binary

Properties of Bomb Phases

**Phases test understanding of different C constructs
and how they are compiled to machine code**

- Phase 1: string comparison
- Phase 2: loop
- Phase 3: switch statement/jump table
- Phase 4: recursive call
- Phase 5: pointers
- Phase 6: linked list/pointers/structs
- Secret phase: binary search (biggest challenge is figuring out how to reach phase)

Phases start out easy and get progressively harder

Let's defuse a bomb phase!

```
0000000000400a6c <phase_2>:  
... # function prologue not shown  
400a72:    mov    %rsp,%rsi  
400a75:    callq  4010ba <read_six_numbers> # rd 6 ints into buffer  
400a7a:    cmpl   $0x1,(%rsp)  
400a7e:    je     400a85 <phase_2+0x19>  
400a80:    callq  400f6d <explode_bomb>  
400a85:    lea    0x4(%rsp),%rbx          # p = &buf[1]  
400a8a:    lea    0x18(%rsp),%rbp         # pend = &buf[6]  
400a8f:    mov    -0x4(%rbx),%eax        # LOOP: v = buf[0]  
400a92:    add    %eax,%eax             # v = 2*v  
400a94:    cmp    %eax,(%rbx)           # if v == *p  
400a96:    je     400a9d <phase_2+0x31> # then goto OK:  
400a98:    callq  400f6d <explode_bomb> # else explode!  
400a9d:    add    $0x4,%rbx            # OK: p++  
400aa1:    cmp    %rbp,%rbx            # if p != pend  
400aa4:    jne    400a8f <phase_2+0x23> # then goto LOOP:  
... # function epilogue not shown  
400aac:    c3                 retq    # YIPPEE!
```

Source Code for Bomb Phase

```
/*
 * phase2b.c - To defeat this stage the user must enter the geometric
 * sequence starting at 1, with a factor of 2 between each number
 */
void phase_2(char *input)
{
    int i;
    int numbers[6];

    read_six_numbers(input, numbers);

    if (numbers[0] != 1)
        explode_bomb();

    for(i = 1; i < 6; i++) {
        if (numbers[i] != numbers[i-1] * 2)
            explode_bomb();
    }
}
```

The Beauty of the Bomb

For the Student

- Get a deep understanding of machine code in the context of a fun game
- Learn about machine code in the context they will encounter in their professional lives
 - Working with compiler-generated code
- Learn concepts and tools of debugging
 - Forward vs backward debugging
 - Students *must* learn to use a debugger to defuse a bomb

For the Instructor

- Self-grading
- Scales to different ability levels
- Easy to generate variants and to port to other machines

Attack Lab

```
int getbuf()
{
    char buf[4];
    /* Read line of text and store in buf */
    gets(buf);
    return 1;
}
```

Task

- Each student assigned “cookie”
 - Randomly generated 8-digit hex string
- Generate string that will cause `getbuf` to return cookie
 - Instead of 1

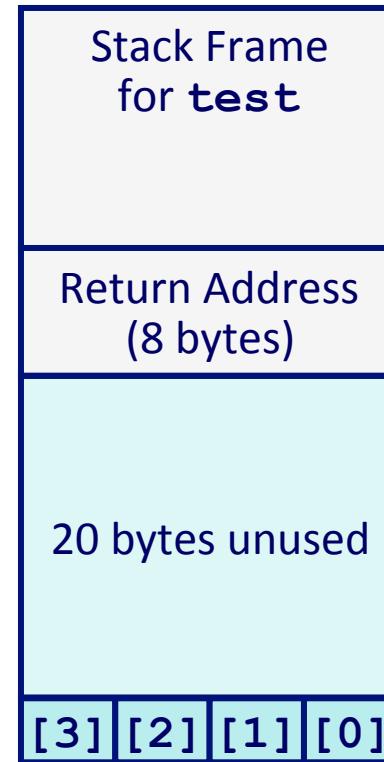
Buffer Code

Return address

```
void test() {  
    int v = getbuf();  
    ...  
}
```

```
void getbuf() {  
    char buf[4];  
    gets(buf);  
    return 1;  
}
```

Stack when gets called



- Calling function `gets(p)` reads characters up to '\n'
- Stores string + terminating null as bytes starting at `p`
- Assumes enough bytes allocated to hold entire string

Buffer Code: Good case

Return address

```
void test() {  
    int v = getbuf();  
    ...  
}
```

```
void getbuf() {  
    char buf[4];  
    gets(buf);  
    return 1;  
}
```

Input string
“01234567890123456789012”

Stack Frame
for **test**

00	00	00	00
00	40	06	f6
00	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

Increasing addresses

buf ← %rsp

- Fits within allocated storage
 - String is 23 characters long + 1 byte terminator

Buffer Code: Bad case

Return address

```
void test() {  
    int v = getbuf();  
    ...  
}
```

```
void getbuf() {  
    char buf[4];  
    gets(buf);  
    return 1;  
}
```

Input string
“0123456789012345678901234”

Stack Frame for test			
00	00	00	00
00	40	00	34
33	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

Increasing addresses

buf ← %rsp

- Overflows allocated storage
 - Corrupts saved frame pointer and return address
- Jumps to address 0x400034 when `getbuf` attempts to return
 - Program executes some instruction and then segfaults

Malicious Use of Buffer Overflow

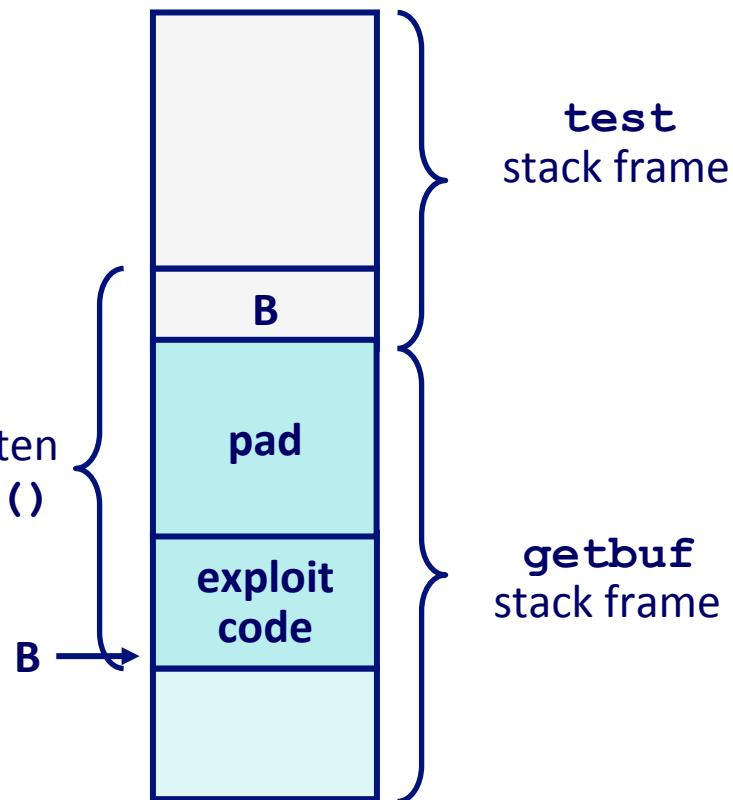
Return address

```
void test() {  
    int v = getbuf();  
    ...  
}
```

```
void getbuf() {  
    char buf[4];  
    gets(buf);  
    return 1;  
}
```

data written
by gets()

Stack after call to `gets()`



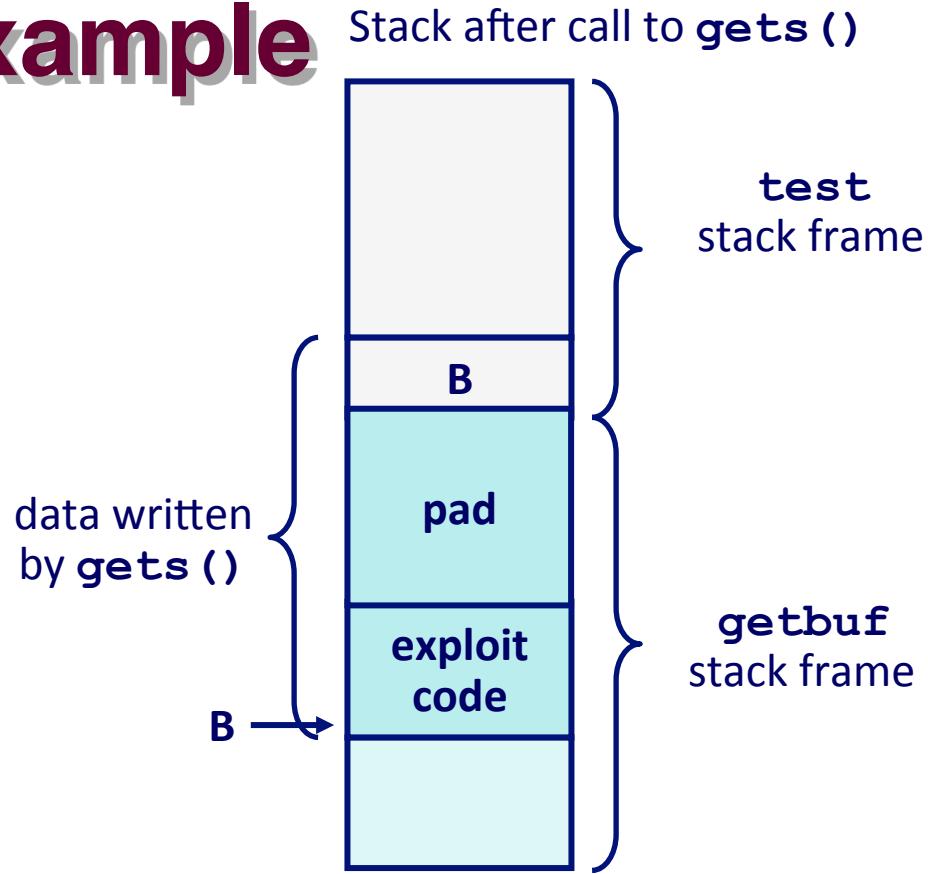
- Input string contains byte representation of executable code
- Overwrite return address with address of buffer
- When `getbuf()` executes return instruction, will jump to exploit code

Exploit String Example

```
void getbuf() {  
    char buf[4];  
    gets(buf);  
    return 1;  
}
```

- Sets 0x59b997fa as function argument
- Invokes function touch2

```
/* Byte code for shell code  
 movq $0x59b997fa,%rdi; ret */  
48 c7 c7 fa 97 b9 59 c3  
/* Pad with 16 bytes */  
00 00 00 00 00 00 00 00  
00 00 00 00 00 00 00 00  
/* Address of shellcode */  
78 dc 61 55 00 00 00 00  
/* Address of touch2 */  
0c 18 40 00 00 00 00 00
```



Why Do We Teach This Stuff?

Important Systems Concepts

- Stack discipline and stack organization
- Instructions are byte sequences
- Making use of tools
 - Debuggers, assemblers, disassemblers

Computer Security

- What makes code vulnerable to buffer overflows
- Common vulnerability in systems

Impact

- CMU student teams consistently win international Capture the Flag Competitions

Cache Lab

Goal: Understanding Cache Operations

- How memory locations map to cache blocks
- Performance implications for application programs

Activities

- Write cache simulator
 - Provides full understanding of mapping from memory address to cache location
- Minimize cache misses for simple application
 - Matrix transpose

Shell Lab

Goal: Write a Unix shell with job control

- (e.g., ctrl-z, ctrl-c, jobs, fg, bg, kill)

Lessons:

- First introduction to systems-level programming and concurrency
- Learn about processes, process control, signals, and catching signals with handlers
- Demystifies command line interface

Infrastructure

- Students use a scripted autograder to incrementally test functionality in their shells

Malloc Lab

Goal: Build your own dynamic storage allocator

```
void *malloc(size_t size)
void *realloc(void *ptr, size_t size)
void free(void *ptr)
```

Lessons

- Sense of programming underlying system
- Large design space with classic time-space tradeoffs
- Develop understanding of scary “action at a distance” property of memory-related errors
- Learn general ideas of resource management

Infrastructure

- Trace driven test harness evaluates implementation for combination of throughput and memory utilization
- Evaluation server and real time posting of scores

Proxy Lab

Goal: write concurrent Web proxy.



Lessons: Ties together many ideas from earlier

- Data representations, byte ordering, memory management, concurrency, processes, threads, synchronization, signals, I/O, network programming, application-level protocols (HTTP)

Infrastructure:

- Plugs directly between existing browsers and Web servers
- Grading is done via autograders and one-on-one demos
- Very exciting for students, great way to end the course

ICS Summary

Principle

- *Introduce students to computer systems from the programmer's perspective rather than the system builder's perspective*

Themes

- What parts of the system affect the correctness, efficiency, and utility of my C programs?
- Makes systems fun and relevant for students
- Prepare students for builder-oriented courses
 - Architecture, compilers, operating systems, networks, distributed systems, databases, ...
 - Since our course provides complementary view of systems, does not just seem like a watered-down version of a more advanced course
 - Gives them better appreciation for what to build

CMU Courses that Build on ICS

CS

Parallel
Prog.

Compilers

Dist.
Systems

Secure
Coding

Networks

Storage
Systems

Software
Engin.

Operating
Systems

Databases

Robotics

Cog.
Robotics

Comp.
Photo.

Computer
Graphics

ECE

Embedded
Control

Real-Time
Systems

Embedded
Systems

Computer
Arch.

ICS

Fostering “Friendly Competition”

Desire

- Challenge the best without frustrating everyone else

Method

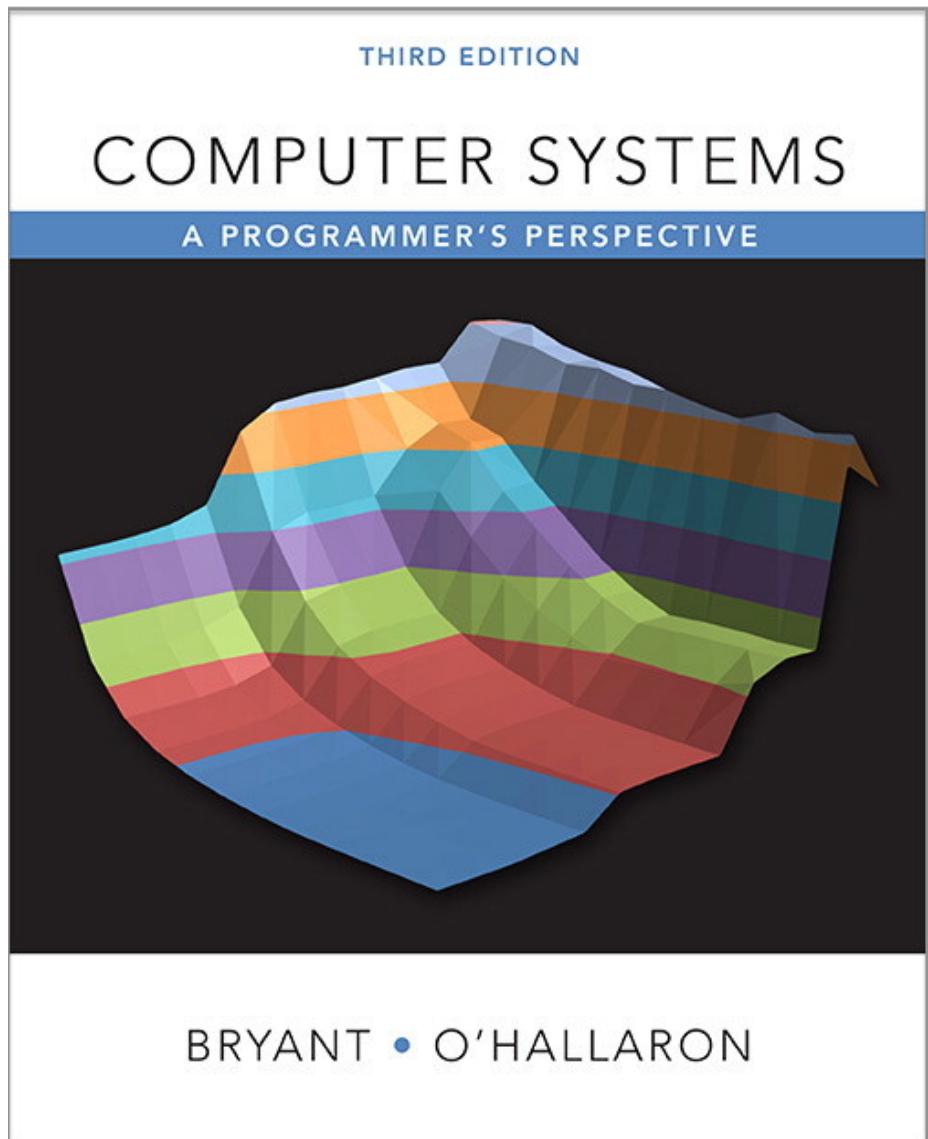
- Web-based submission of solutions
- Server checks for correctness and computes performance score
 - How many stages passed, program throughput, ...
- Keep updated results on web page
 - Students choose own *nom de guerre*

Relationship to Grading

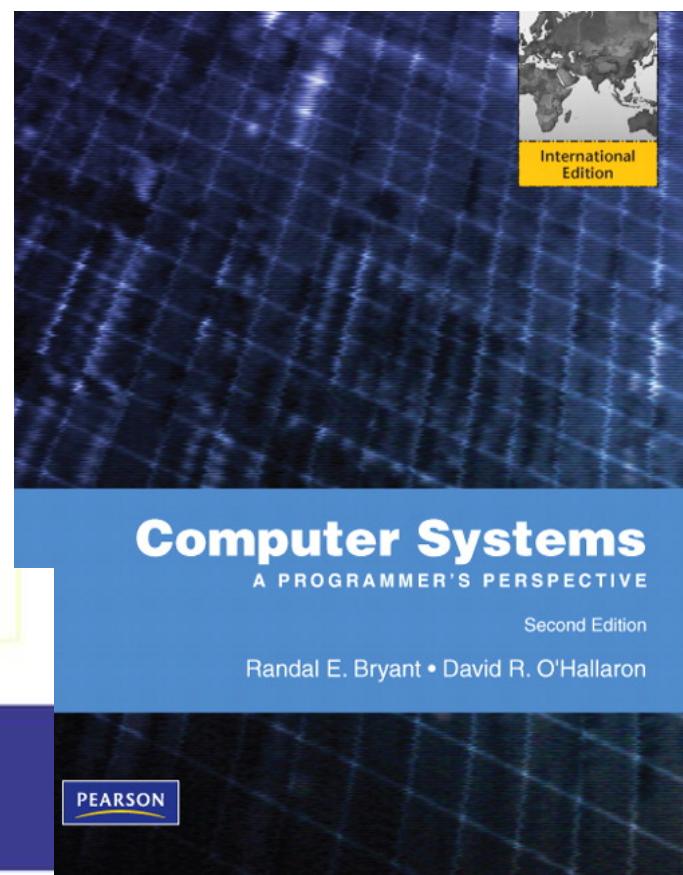
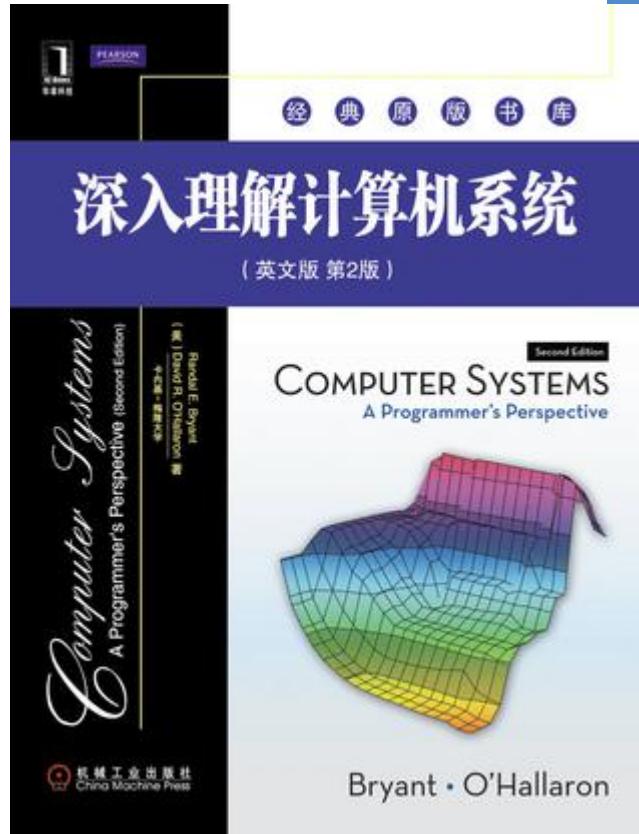
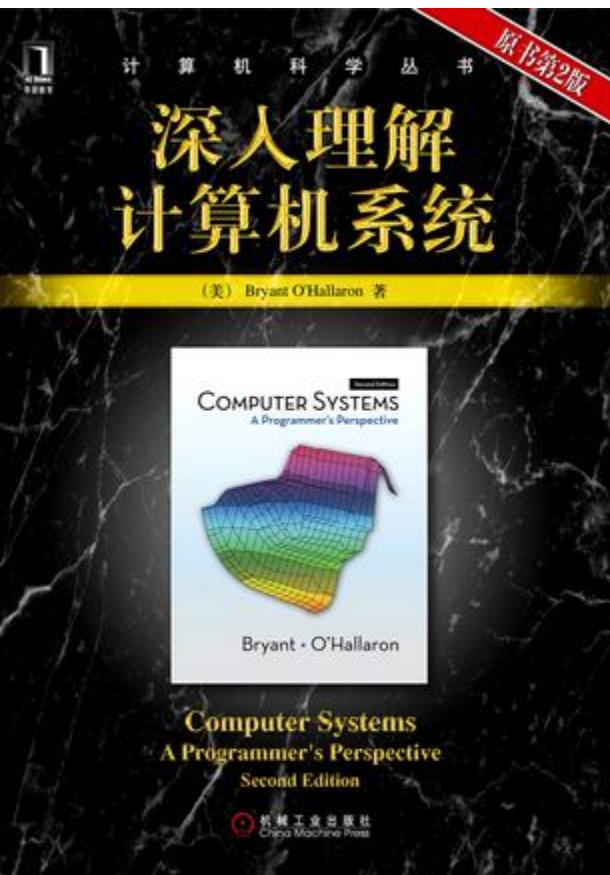
- Students get full credit once they reach set threshold
- Push beyond this just for own glory/excitement

Shameless Promotion

- <http://csapp.cs.cmu.edu>
- Third edition published 2015
- In use at 289 institutions worldwide

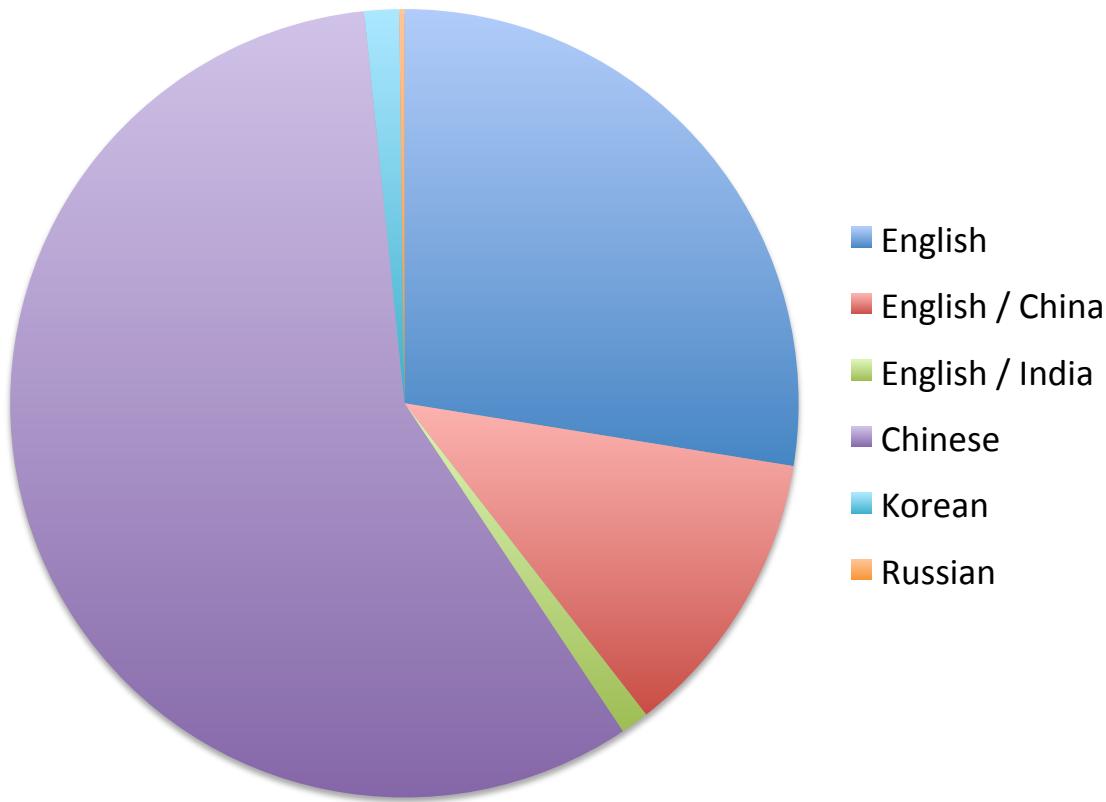


International Editions (No 3rd edition yet)



Overall Sales

- All Editions
- As of 6/30/2015
- 175,835 total

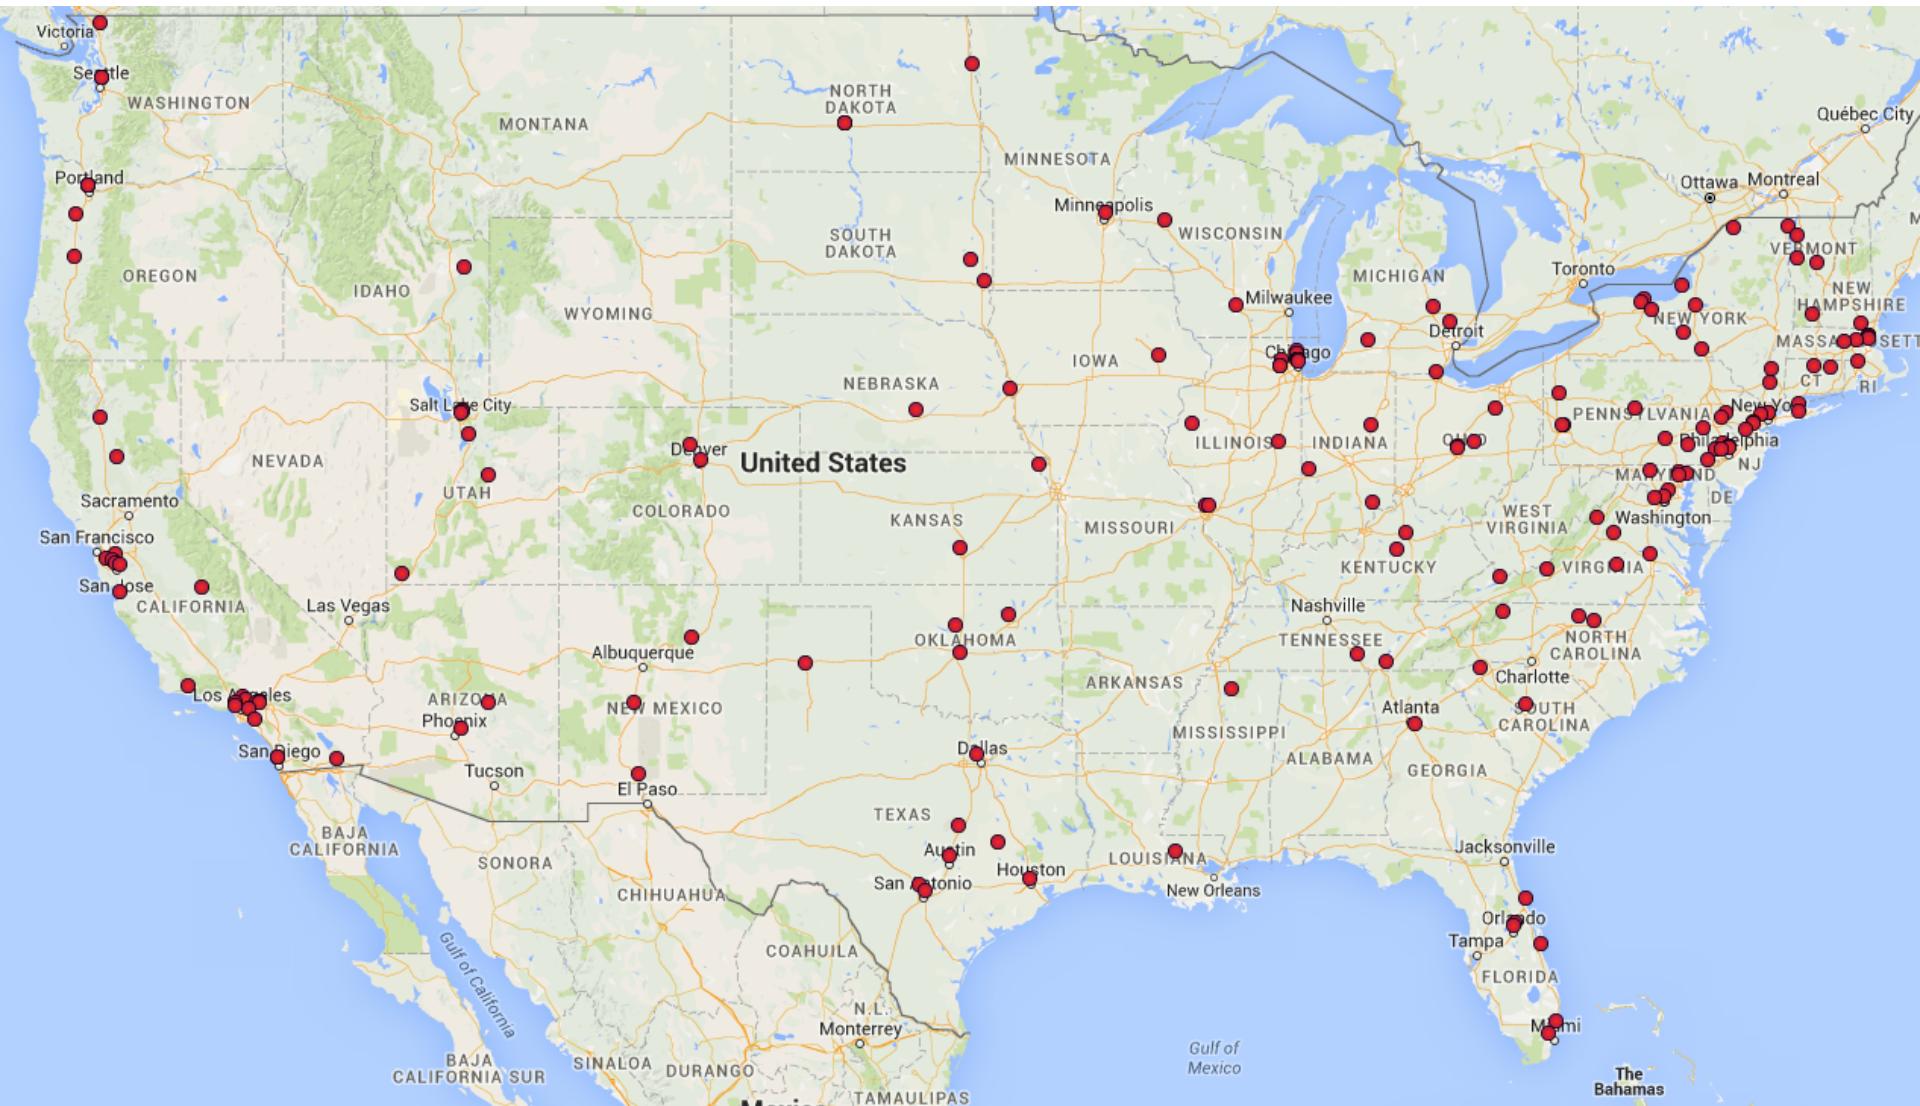


Worldwide Adoptions



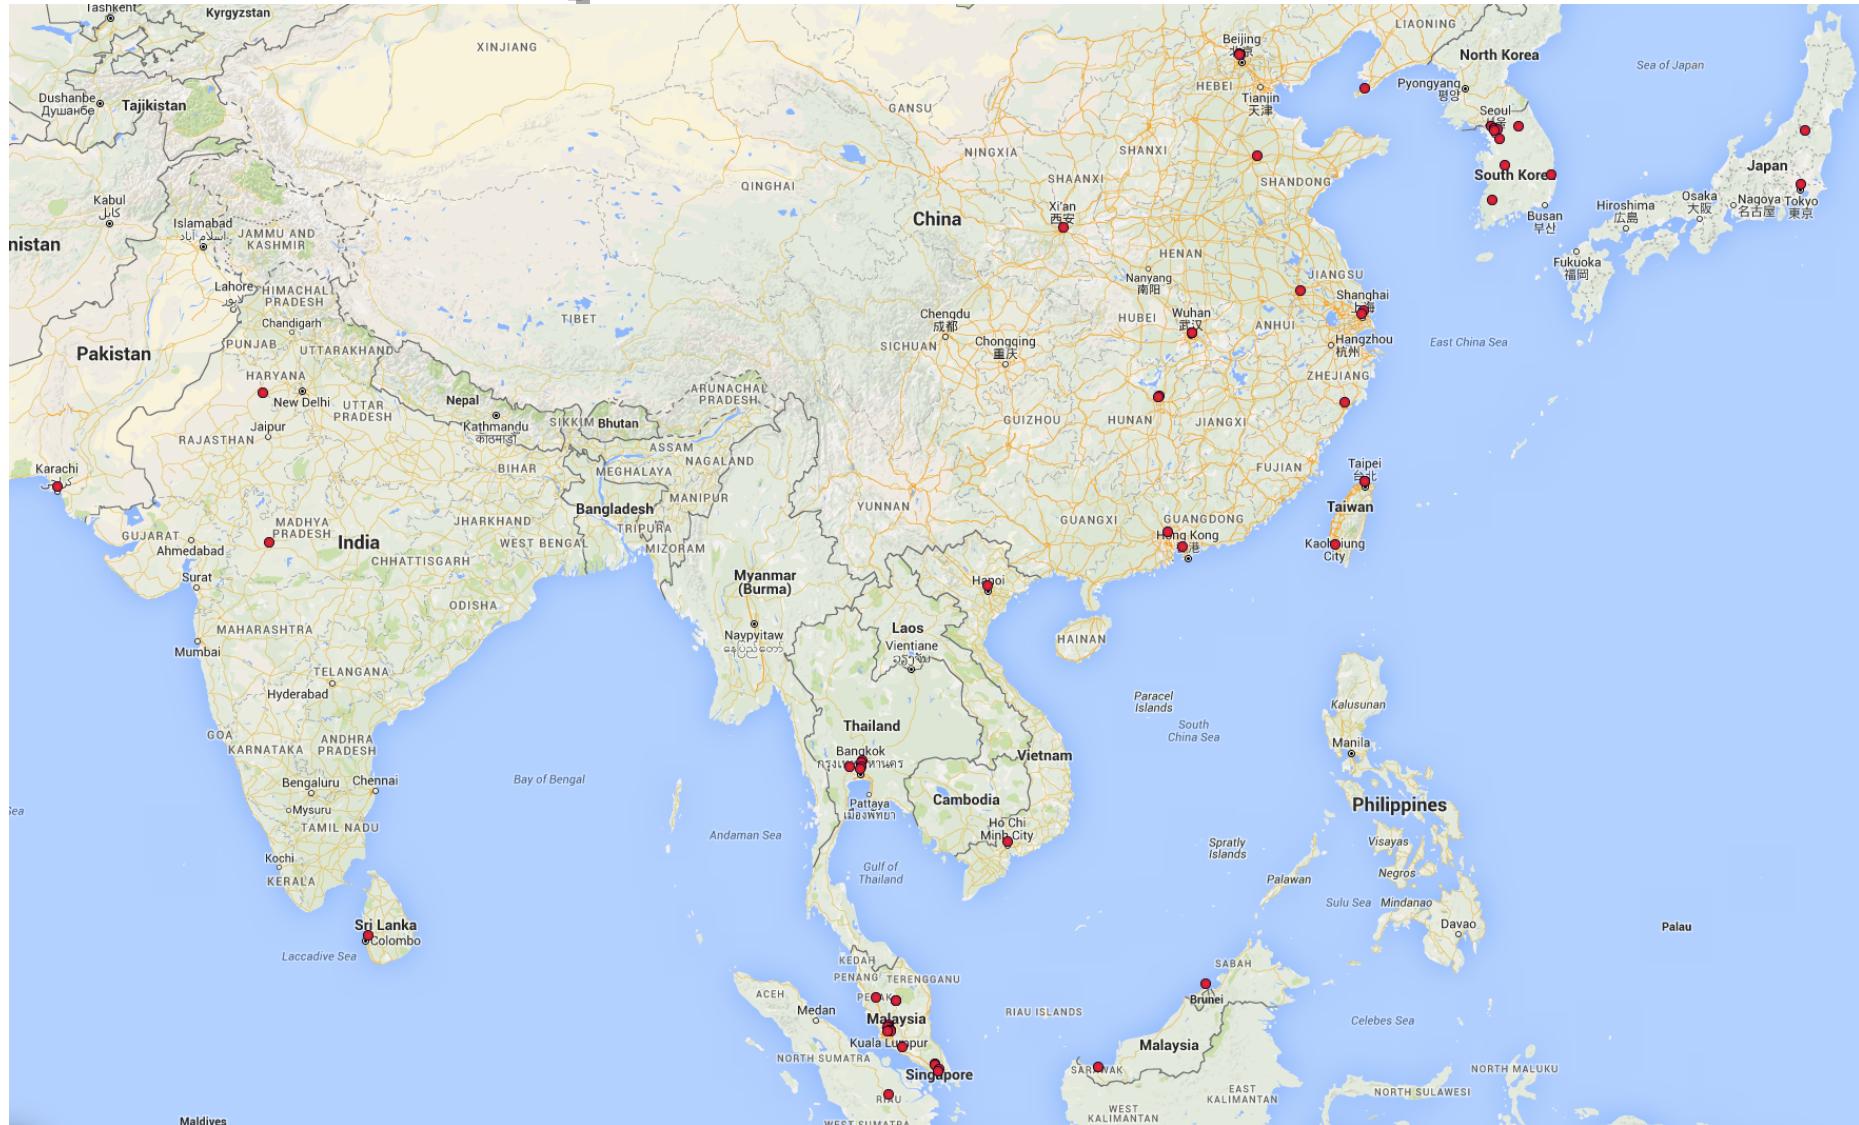
289 total

US Adoptions



176 total

Asian Adoptions



European Adoptions



CS:APP3e

Vital stats:

- 12 chapters
- 267 practice problems (solutions in book)
- 226 homework problems (solutions in instructor's manual)
- 544 figures, 342 line drawings
- Many C & machine code examples

Turn-key course provided with book:

- Electronic versions of all code examples.
- Powerpoint and PDF versions of each line drawing
- Password-protected Instructors Page
 - Instructor's Manual
 - Lab Infrastructure
 - Powerpoint lecture notes
 - Exam problems.

Coverage

Material Used by ICS at CMU

- Pulls together material previously covered by multiple textbooks, system programming references, and man pages

Greater Depth on Some Topics

- Dynamic linking
- I/O multiplexing

Additional Topic

- Computer Architecture
- Added to cover all topics in “Computer Organization” course

Architecture

Material

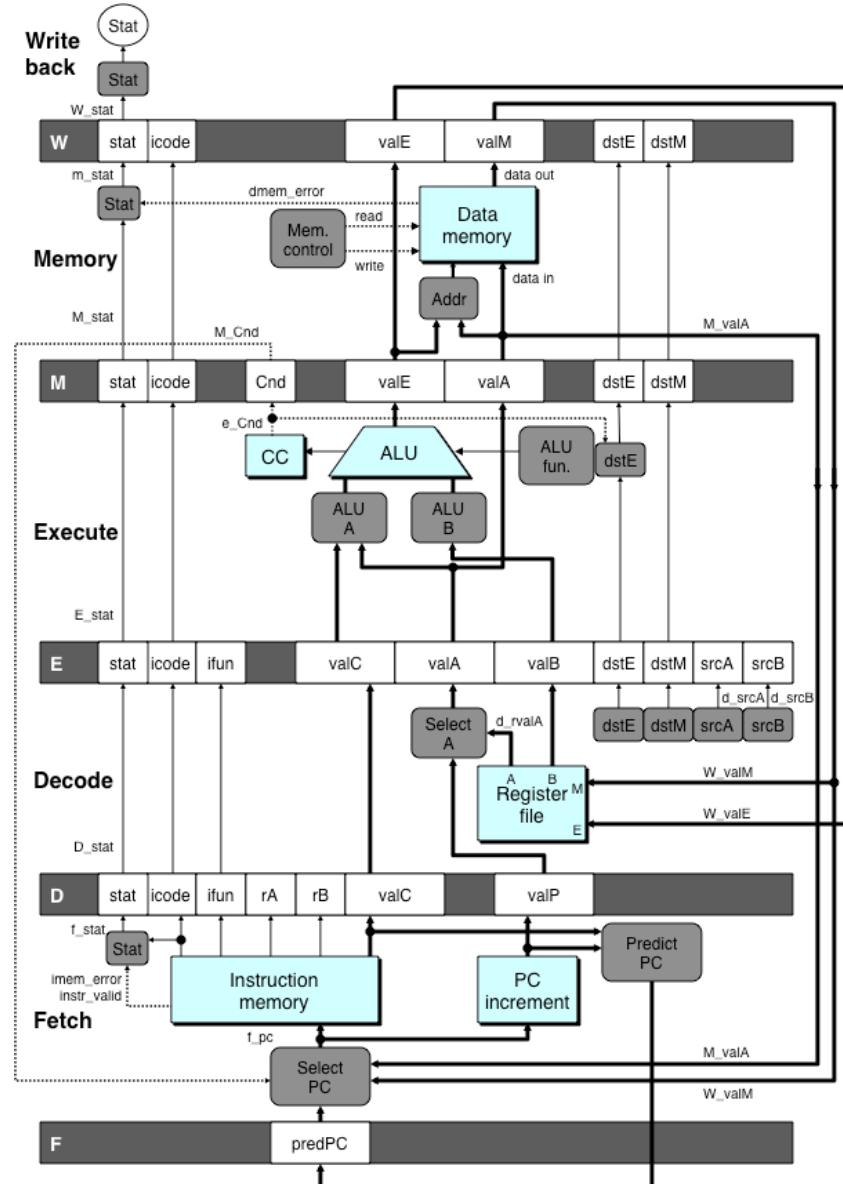
- Y86-64 instruction set
 - Simplified/reduced x86-64
- Implementations
 - Sequential
 - 5-stage pipeline

Presentation

- Simple hardware description language to describe control logic
- Automatic translation to simulator and to Verilog

Labs

- Modify / extend processor design
 - New instructions
 - Change branch prediction policy
- Optimize application + processor



Web Asides

- Supplementary material via web
- Topics either more advanced or more arcane

Examples

- Boolean algebra & Boolean rings
- IA32 programming
- Combining assembly & C code
- Processor design in Verilog
- Using SIMD instructions
- Memory blocking

Courses Based on CS:APP

Computer Organization

- ORG** Topics in conventional computer organization course, but with a different flavor
- ORG+** Extends computer organization to provide more emphasis on helping students become better application programmers

Introduction to Computer Systems

- ICS** Create enlightened programmers who understand enough about processor/OS/compilers to be effective
- ICS+** What we teach at CMU. More coverage of systems software

Systems Programming

- SP** Prepare students to become competent system programmers

Courses Based on CS:APP

Chapter	Topic	Course				
		ORG	ORG+	ICS	ICS+	SP
1	Introduction	●	●	●	●	●
2	Data representations	●	●	●	●	○
3	Machine language	●	●	●	●	●
4	Processor architecture	●	●			
5	Code optimization		●	●	●	
6	Memory hierarchy	○	●	●	●	○
7	Linking			○	○	●
8	Exceptional control flow			●	●	●
9	Virtual memory	○	●	●	●	●
10	System-level I/O				●	●
11	Concurrent programming				●	●
12	Network programming				●	●

○ Partial Coverage

● Complete Coverage

Conclusions

ICS Has Proved Its Success

- Thousands of students at CMU over 13 years
- Positive feedback from alumni
- Positive feedback from systems course instructors

CS:APP is International Success

- Supports variety of course styles
- Many purchases for self study